Week 5: Optimisation

Made-up lecture for tests

Gradient descent

- Step against the gradient
- The learning rate sets the step size
- Too large a rate makes the loss bounce between the walls of the valley and sometimes diverge entirely

Momentum

- Keep a running average of past gradients
- Heavy ball: $\beta = 0.9$ is common